ECON-613 HW #2

Min Chul Kim

2019-02-15

Exercise1

```r
#Set Seed
set.seed(1)

#Drawing Samples from Distributions
n = 10000

X1 = runif(n, min = 1, max = 3)
X2 = rgamma(n, shape = 3, scale = 2)
X3 = rbinom(n, size = 1, prob = 0.3)
eps = rnorm(n, mean = 2, sd = 1)

#Creating Y and Dummy Y Variables
Y = 0.5 + 1.2*X1 - 0.9*X2 + 0.1*X3 + eps
ydum = ifelse(Y > mean(Y), 1, 0)
```

Exercise2

```r
#Calculating the Correlation between Y and X1, How Different is it from 1.2
correlation = cor(Y, X1)
difference_corr = abs(cor(Y, X1) - 1.2) %>% as.data.frame
colnames(difference_corr) = "Difference between Correlation and 1.2"

kable(difference_corr, digits = 4, format = "markdown")
```

Difference between Correlation and 1.2
0.9837

```r
#Regression
X = data.frame(
  1,
  X1,
  X2,
  X3
) %>%
  as.matrix()

#Calculating the Coefficients of this Regression
reg1_coef = solve((t(X) %*% X)) %*% t(X) %*% Y

df_reg1_coef = reg1_coef %>% as.data.frame
```

```r
colnames(df_reg1_coef) = "Coefficients of Reg"
rownames(df_reg1_coef) = c("Int", "X1", "X2", "X3")

kable(df_reg1_coef, digits = 4,
      format = "markdown", caption = "Coefficients of Regression")
```

	Coefficients of Reg
Int	2.4919
X1	1.1965
X2	-0.8978
X3	0.1251

```r
#Caluclating the Standard Errors Using Standard Formula
k = length(reg1_coef)
est_var = sum( (Y - (X %*% reg1_coef))^2 ) / (n - k)
est_std = est_var * solve( (t(X) %*% X) ) %>%
  diag() %>%
  sqrt()

df_est_std = est_std %>% as.data.frame
colnames(df_est_std) = "SE from Standard Formula"
rownames(df_est_std) = c("Int", "X1", "X2", "X3")

kable(df_est_std, digits = 4,
      format = "markdown", caption = "SE from Standard Formula")
```

	SE from Standard Formula
Int	0.0410
X1	0.0174
X2	0.0029
X3	0.0221

```r
#Bootstrap n = 49
boot_coef = list()

for(i in 1:49){

  boot_x1 = sample(X1, n, replace = TRUE)
  boot_x2 = sample(X2, n, replace = TRUE)
  boot_x3 = sample(X3, n, replace = TRUE)

  boot_X = cbind(1, boot_x1, boot_x2, boot_x3)
  boot_Y = sample(Y, n, replace = TRUE)

  boot_coef[[i]] = solve((t(boot_X) %*% boot_X)) %*% t(boot_X) %*% boot_Y
}

#Making Bootstrap Values into a Dataframe
boot_temp_df = map(boot_coef, data.frame)
boot_coef_df = do.call("cbind", boot_temp_df)
```

```
colnames(boot_coef_df) = paste("sample", 1:49, sep = "")

#Calculating SE
bootstrap_se1 = apply(boot_coef_df, 1, FUN = sd) %>% as.data.frame()
colnames(bootstrap_se1) = "SE of Bootstrap 49"
rownames(bootstrap_se1) = c("Int", "X1", "X2", "X3")
kable(bootstrap_se1, digits = 4, format = "markdown")
```

	SE of Bootstrap 49
Int	0.1227
X1	0.0477
X2	0.0100
X3	0.0656

```
#Bootstrap n = 499
boot_coef = list()

for(i in 1:499){

  boot_x1 = sample(X1, n, replace = TRUE)
  boot_x2 = sample(X2, n, replace = TRUE)
  boot_x3 = sample(X3, n, replace = TRUE)

  boot_X = cbind(1, boot_x1, boot_x2, boot_x3)
  boot_Y = sample(Y, n, replace = TRUE)

  boot_coef[[i]] = solve((t(boot_X) %*% boot_X)) %*% t(boot_X) %*% boot_Y
}

#Making Bootstrap Values into a Dataframe
boot_temp_df = map(boot_coef, data.frame)
boot_coef_df = do.call("cbind", boot_temp_df)
colnames(boot_coef_df) = paste("sample", 1:499, sep = "")

#Calculating SE
bootstrap_se1 = apply(boot_coef_df, 1, FUN = sd) %>% as.data.frame()
colnames(bootstrap_se1) = "SE of Bootstrap 499"
rownames(bootstrap_se1) = c("Int", "X1", "X2", "X3")
kable(bootstrap_se1, digits = 4, format = "markdown")
```

	SE of Bootstrap 499
Int	0.1461
X1	0.0610
X2	0.0094
X3	0.0721

Answer:
The difference in correlation and 1.2 is about 1. Personally judged, this difference is significant.

Exercise3

In the following R chunk, I create an ordinary probit likelihood function, as well as the probit log likelihood function. For the ordinary likelihood function, expect 0 for the answer, since probabilities less than or equal to 1 are being multiplied for a number of times. For the log likelihood function, expect a positive result, because I am supplying the negative of the computed log likelihood.

```r
#Creating Likelihood Function
probit_likelihood = function(x, beta, ydum){

  #Calculating XB
  X_Beta = x %*% beta

  #Calculating Phi(XB)
  phi = map(X_Beta, pnorm) %>% unlist()

  #Calculating the Likelihood Using Likelihood Formula
  prob = phi^(ydum) * (1 - phi)^(1 - ydum)
  likelihood = prod( prob )

  return(likelihood)
}


#The following will be 0 since probabilities <= 1 are being multiplied.
pro_df = data.frame(
  probit_likelihood(X, reg1_coef, ydum)
)
colnames(pro_df) = "Probit Likelihood"

kable(pro_df, digits = 4, format = "markdown")
```

Probit Likelihood
0

```r
#Creating Log Likelihood Function
pro_log_likelihood = function(x, beta, ydum){

  #Calculating XB
  X_Beta = x %*% beta

  #Calculating Phi(XB) and Preventing the Obtained Results from Being 0 and 1
  phi = map(X_Beta, pnorm) %>% unlist()
  phi[phi == 0] = 0.0001
  phi[phi == 1] = 0.9999

  #Calculating the Log Likelihood Using the Formula
  log_likelihood =  sum( ydum * log(phi) + (1 - ydum) * log((1 - phi)) )

  return(-log_likelihood)
}

pro_log_df = data.frame(
  pro_log_likelihood(X, reg1_coef, ydum)
```

```
)
colnames(pro_log_df) = "Probit Log Likelihood"

kable(pro_log_df, digits = 4, format = "markdown")
```

	Probit Log Likelihood
	2506.252

```
#Steepest Ascent Function
steepest_ascent =
  function(x, betas_old, h, ydum_, learning_rate, epsilon){

  #Initial Values (betas_new is initialized to start the while loop)
  d = 10
  betas_new = rep(0, 4)

  while( d > epsilon ){

    #Computing Matrices to Be Used in Derivative
    betas_old_matrix = matrix(betas_old, nrow = 4, ncol = 4)
    betas_new_matrix = betas_old_matrix + h * diag(1, nrow = 4, ncol = 4)

    #Calculating Partial Derivatives of Log Likelihood
    derivative = c()
    for(i in 1:4){
      derivative[i] = ( pro_log_likelihood(x, betas_new_matrix[,i], ydum_) -
                        pro_log_likelihood(x, betas_old_matrix[,i], ydum_) ) / h
    }

    #Updating Betas
    betas_old = betas_new
    betas_new = betas_old - learning_rate * derivative

    #Calculating sum_betas_diff
    d = sum( (betas_new - betas_old)^2 )

  }

  return(betas_new)
}

#Implementing Steepest Ascent
betas = rep(0, 4)

steepas = steepest_ascent(X, betas_old = betas,
              h = 10^(-6), ydum, learning_rate = 0.000001, epsilon = .00001)

#True Parameters
probit_model = glm(ydum ~ X1 + X2 + X3, family = binomial(link = "probit"))

## Warning: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
probit_coef = probit_model$coefficients

#Comparison
probit_steepest_df = data.frame(
  steepas,
  probit_coef,
  abs(steepas - probit_coef)
)
colnames(probit_steepest_df) = c("Steepest Ascent", "True", "Difference")

kable(probit_steepest_df, digits = 4, format = "markdown")
```

	Steepest Ascent	True	Difference
(Intercept)	0.2059	3.0846	2.8787
X1	0.4462	1.1296	0.6834
X2	-0.2013	-0.8885	0.6872
X3	0.0660	0.1830	0.1170

Answer:
According to the result, my steepest ascent values are fairly comparable to the true parameter quantities. The sign of each parameter matches, as can be seen in the table. The variable that deviates the most from the true parameter is my intercept, while the variable that deviates the least from the true parameter is my X3.

Exercise 4

```
##Writing and Optimizing Linear Probability Model
linear_reg = lm(ydum ~ X1 + X2 + X3)

dat = data.frame(
  X,
  ydum
)
colnames(dat) = c("Int", "X1", "X2", "X3", "Y")

#RSS to Minimize Using Optim Function
min_RSS = function(data, par){
  with(data, sum( (par[1] + par[2] * X1 + par[3] * X2 + par[4] * X3 - ydum)^2 ))
}

linear_coef_df = data.frame(
  optim(par = c(0, 1, 1, 1), min_RSS, data = dat)$par,
  linear_reg$coefficients
)
colnames(linear_coef_df) = c("Optim Linear Function", "lm Linear Function")

kable(linear_coef_df, digits = 4, format = "markdown")
```

	Optim Linear Function	lm Linear Function
(Intercept)	0.9108	0.9133

	Optim Linear Function	lm Linear Function
X1	0.1315	0.1312
X2	-0.1021	-0.1023
X3	0.0231	0.0232

```r
##Writing and Optimizing Logit Model
logit_reg = glm(ydum ~ X1 + X2 + X3, family = binomial(link = "logit"))

sigmoid = function(x){
  return( 1 / (1 + exp(-x)) )
}

logit_cost_function = function(par){

  len = length(X)
  g = sigmoid(X %*% par)
  C = -(1 / len) * sum( ydum * log(g) + (1 - ydum) * log(1 - g))
  return(C)

}

initial_par = rep(0, ncol(X))

logit_coef_df = data.frame(
  optim(par = initial_par, fn = logit_cost_function)$par,
  logit_reg$coefficients
)
colnames(logit_coef_df) = c("Optim logit Function", "glm logit Function")

kable(logit_coef_df, digits = 4, format = "markdown")
```

	Optim logit Function	glm logit Function
(Intercept)	5.5030	5.5030
X1	2.0297	2.0299
X2	-1.5910	-1.5912
X3	0.3277	0.3281

```r
##Writing and Optimizing Probit Model
probit_reg = glm(ydum ~ X1 + X2 + X3, family = binomial(link = "probit"))

probit_cost_function = function(par){

  len = length(X)
  phi = pnorm(X %*% par)
  C = -(1 / len) * sum ( ydum * log(phi) + (1 - ydum) * log(1 - phi) )
  return(C)

}

probit_coef_df = data.frame(
optim(par = initial_par, fn = probit_cost_function)$par,
```

```
probit_reg$coefficients
)
colnames(probit_coef_df) = c("Optim Probit Function", "glm Probit Function")

kable(probit_coef_df, digits = 4, format = "markdown")
```

	Optim Probit Function	glm Probit Function
(Intercept)	3.0839	3.0846
X1	1.1296	1.1296
X2	-0.8884	-0.8885
X3	0.1837	0.1830

Interpretation:

Linear Regression:

Values from the optim function and the base package lm function are not different at all. These differences can be considered neglegible.

Intercept: When all the X's are 0, the expected ydum variable is 0.91.
X1: Ceteris paribus, an increase in X1 increases the expected ydum value.
X2: Ceteris paribus, an increase in X2 decreases the expected ydum value.
X3: Ceteris paribus, an increase in X3 increases the expected ydum value.

Logistic Regression:

Values from the optim function and the base package glm function are not different at all. These differences can be considered neglegible.

Intercept: When all the X's are 0, the expected log odds of ydum = 1 is approximately 5%
X1: Ceteris parbius, an increase in X1 increases the likelihood of ydum = 1.
X2: Ceteris paribus, an increase in X2 decreases the likelihood of ydum = 1.
X3: Ceteris paribus, an increase in X3 increases the likelihood of ydum = 1.

Probit Regression:

Values from the optim function and the base package glm function are not different at all. These differences can be considered neglegible.

Intercept: When all the X's are 0, the expected likelihood ydum = 1 under the Standard Normal distribution is approximately 3.1 %
X1: Ceteris paribus, an increase in X1 increases the likelihood of ydum = 1.
X2: Ceteris paribus, an increase in X2 decreases the likelihood of ydum = 1.
X3: Ceteris paribus, an increase in X3 increases the likelihood of ydum = 1.

Exercise 5

```r
#Probit Model Marginal Effects
probit_x_beta = predict(probit_reg, type = "link")
probit_pdf = dnorm(probit_x_beta) %>% mean()
probit_marginal = probit_pdf * probit_reg$coefficients

#logit Model Marginal Effects
logit_x_beta = predict(logit_reg, type = "link")
logit_pdf = dlogis(logit_x_beta) %>% mean()
logit_marginal = logit_pdf * logit_reg$coefficients

#Summary Table
marginals_df = data.frame(
  probit_marginal,
  logit_marginal
)
colnames(marginals_df) = c("Probit", "Logit")

kable(marginals_df, digits = 4, format = "markdown")
```

	Probit	Logit
(Intercept)	0.3741	0.3726
X1	0.1370	0.1374
X2	-0.1078	-0.1077
X3	0.0222	0.0222

```r
#Bootstrap for Probit
boot_probit_coef = list()
for(i in 1:49){

  boot_probit_x_beta = sample(probit_x_beta, n, replace = TRUE)
  boot_probit_pdf = dnorm(boot_probit_x_beta) %>% mean()
  boot_probit_marginal = boot_probit_pdf * probit_reg$coefficients

  boot_probit_coef[[i]] = boot_probit_marginal

}

#Bootstrap for Logit
boot_logit_coef = list()
for(i in 1:49){

  boot_logit_x_beta = sample(logit_x_beta, n, replace = TRUE)
  boot_logit_pdf = dnorm(boot_logit_x_beta) %>% mean()
  boot_logit_marginal = boot_logit_pdf * logit_reg$coefficients

  boot_logit_coef[[i]] = boot_logit_marginal

}

#Making Bootstrap Values into a Dataframe
```

```r
boot_probit_temp_df = map(boot_probit_coef, data.frame)
boot_probit_coef_df = do.call("cbind", boot_probit_temp_df)
colnames(boot_probit_coef_df) = paste("sample", 1:49, sep = "")

boot_logit_temp_df = map(boot_logit_coef, data.frame)
boot_logit_coef_df = do.call("cbind", boot_logit_temp_df)
colnames(boot_logit_coef_df) = paste("sample", 1:49, sep = "")

#Calculating SD for Probit
bootstrap_probit_sd = apply(boot_probit_coef_df, 1, FUN = sd) %>% as.data.frame()
colnames(bootstrap_probit_sd) = "Probit SD of Bootstrap 49"
rownames(bootstrap_probit_sd) = c("Int", "X1", "X2", "X3")
kable(bootstrap_probit_sd, digits = 4, format = "markdown")
```

	Probit SD of Bootstrap 49
Int	0.0051
X1	0.0019
X2	0.0015
X3	0.0003

```r
#Calculating SD for Logit
bootstrap_logit_sd = apply(boot_logit_coef_df, 1, FUN = sd) %>% as.data.frame()
colnames(bootstrap_logit_sd) = "Logit SD of Bootstrap 49"
rownames(bootstrap_logit_sd) = c("Int", "X1", "X2", "X3")
kable(bootstrap_logit_sd, digits = 4, format = "markdown")
```

	Logit SD of Bootstrap 49
Int	0.0061
X1	0.0023
X2	0.0018
X3	0.0004

```r
#Delta Method

##Computing Covariance Matrix of Probit Regression
v = vcov(probit_reg)

##Function to be used for Gradient
g = function(x){

  probit_pdf = dnorm(X %*% x) %>% mean()
  probit_marginal = probit_pdf * x

  return(probit_marginal)
}

##Gradient
gradient_g = numDeriv::jacobian(g, probit_reg$coefficients)

##Table of Delta Method SD
```

```
Delta_df = gradient_g %*%
  v %*%
  t(gradient_g) %>%
  diag() %>%
  sqrt() %>%
  as.data.frame()
colnames(Delta_df) = "Delta Method SD"
rownames(Delta_df) = c("Int", "X1", "X2", "X3")

kable(Delta_df, digits = 4, format = "markdown")
```

	Delta Method SD
Int	0.0094
X1	0.0044
X2	0.0004
X3	0.0057